

# Monte Carlo methods

Nordic Probabilistic AI School

**Instructor:**

Charles Margossian (Flatiron Institute, New York, USA)

**Teaching Assistants:**

Austin Garrett (?)

Bob Pepin (University of Copenhagen, Denmark)

## Outline:

- Aim and scope of Monte Carlo methods
- Markov chain Monte Carlo
- Application: Bayesian linear regression
- Application: Disease transmission model
- Importance sampling and model comparison
- Discussion

**Remark:** This is a course on effectively applying Monte Carlo methods, rather than developing them from scratch.

However, principled use of MC methods requires a rigorous understanding of how they are implemented. So we need to talk about theory and low-level implementation—all whilst remaining pragmatic.

## Code for exercises:

- We'll use R scripts and the probabilistic programming language **Stan** .
- This is not a course on coding in R: the methods apply generally and **Stan** can be interfaced with other languages, including Python and Julia.
- Code for the exercises at [https://github.com/charlesm93/stanTutorial/tree/main/Nordic\\_Prob\\_AI](https://github.com/charlesm93/stanTutorial/tree/main/Nordic_Prob_AI)

# I

Aim and scope of Monte Carlo methods

**Goal:** compute an intractable expectation value with respect to  $\pi(\theta)$ .

**Goal:** compute an intractable expectation value with respect to  $\pi(\theta)$ .

- Statistical physics

$$\pi(\theta) \propto \exp(-E(\theta))$$

- Bayesian inference

$$\pi(\theta) = p(\theta | y) \propto p(\theta)p(y | \theta)$$

- Variational inference

$$\pi(\theta) = q(\theta)$$

**Goal:** compute an intractable expectation value with respect to  $\pi(\theta)$ .

- Statistical physics

$$\pi(\theta) \propto \exp(-E(\theta))$$

- Bayesian inference

$$\pi(\theta) = p(\theta | y) \propto p(\theta)p(y | \theta)$$

- Variational inference

$$\pi(\theta) = q(\theta)$$

What might we want to learn about these distributions?



**Goal:** compute an intractable expectation value with respect to  $\pi(\theta)$ .

- Statistical physics

$$\pi(\theta) \propto \exp(-E(\theta))$$

- Bayesian inference

$$\pi(\theta) = p(\theta | y) \propto p(\theta)p(y | \theta)$$

- Variational inference

$$\pi(\theta) = q(\theta)$$

What might we want to learn about these distributions?

- Expectation, variance, and quantiles of  $f(\theta)$  with respect to  $\pi$ .

**Goal:** compute an intractable expectation value with respect to  $\pi(\theta)$ .

- Statistical physics

$$\pi(\theta) \propto \exp(-E(\theta))$$

- Bayesian inference

$$\pi(\theta) = p(\theta | y) \propto p(\theta)p(y | \theta)$$

- Variational inference

$$\pi(\theta) = q(\theta)$$

What might we want to learn about these distributions?

- Expectation, variance, and quantiles of  $f(\theta)$  with respect to  $\pi$ .
- **Monte Carlo:** draw samples and construct sample estimators,

$$\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(N)} \sim \pi(\theta).$$

- When no exact simulation is possible, use Markov chain Monte Carlo, or importance sampling.

# II

Markov chain Monte Carlo

Quantities of interest can often be expressed as integrals with respect to a probability measure

$$\mathbb{E}[f(\theta)] = \int f(\theta) p(\theta \mid y) \, \mathrm{d}\theta$$

Quantities of interest can often be expressed as integrals with respect to a probability measure

$$\mathbb{E}[f(\theta)] = \int f(\theta) p(\theta \mid y) \, \mathrm{d}\theta$$

Monte Carlo estimator:

$$\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(N)} \stackrel{\text{iid}}{\sim} p(\theta \mid y)$$

$$\widehat{\mathbb{E}}[f(\theta)] = \frac{1}{N} \sum_{n=1}^N f\left(\theta^{(n)}\right)$$

Quantities of interest can often be expressed as integrals with respect to a probability measure

$$\mathbb{E}[f(\theta)] = \int f(\theta) p(\theta \mid y) \, \mathrm{d}\theta$$

Monte Carlo estimator:

$$\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(N)} \stackrel{\text{iid}}{\sim} p(\theta \mid y)$$

$$\widehat{\mathbb{E}}[f(\theta)] = \frac{1}{N} \sum_{n=1}^N f\left(\theta^{(n)}\right)$$

Can get a sample estimator for mean, variance and quantiles.

How good is our Monte Carlo estimator  $\hat{\mathbb{E}}[f(\theta)]$ ?

How good is our Monte Carlo estimator  $\widehat{\mathbb{E}}[f(\theta)]$ ?

Ultimately want to control the expected squared error,

$$\mathbb{E} \left[ \left( \widehat{\mathbb{E}}[f(\theta)] - \mathbb{E}[f(\theta)] \right)^2 \right] = \text{Bias}^2 + \text{Var} \left[ \widehat{\mathbb{E}}[f(\theta)] \right]$$



How good is our Monte Carlo estimator  $\widehat{\mathbb{E}}[f(\theta)]$ ?

Ultimately want to control the expected squared error,

$$\mathbb{E} \left[ \left( \widehat{\mathbb{E}}[f(\theta)] - \mathbb{E}[f(\theta)] \right)^2 \right] = \text{Bias}^2 + \text{Var} \left[ \widehat{\mathbb{E}}[f(\theta)] \right]$$

If  $\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(N)}$  are i.i.d,

$$\text{Bias} = 0, \quad \text{Var} \left[ \widehat{\mathbb{E}}[f(\theta)] \right] = \frac{1}{N} \text{Var}[\theta]$$

How good is our Monte Carlo estimator  $\widehat{\mathbb{E}}[f(\theta)]$ ?

Ultimately want to control the expected squared error,

$$\mathbb{E} \left[ \left( \widehat{\mathbb{E}}[f(\theta)] - \mathbb{E}[f(\theta)] \right)^2 \right] = \text{Bias}^2 + \text{Var} \left[ \widehat{\mathbb{E}}[f(\theta)] \right]$$

If  $\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(N)}$  are i.i.d,

$$\text{Bias} = 0, \quad \text{Var} \left[ \widehat{\mathbb{E}}[f(\theta)] \right] = \frac{1}{N} \text{Var}[\theta]$$

We also have a *central limit theorem*, i.e. for large  $N$

$$\widehat{\mathbb{E}}[f(\theta)] \overset{\text{approx}}{\sim} \text{normal} \left( \mathbb{E}f(\theta), \sqrt{\frac{\text{Var}[f(\theta)]}{N}} \right).$$

In practice, we cannot generate iid samples from  $p(\theta \mid y)$ .

In practice, we cannot generate iid samples from  $p(\theta \mid y)$ .

Markov chain Monte Carlo:

- Start with an initial draw  $\theta^{(0)} \sim p_0(\theta)$ .
- Apply a transition kernel,  $\theta^{(i+1)} \sim \Gamma(\theta^{(i+1)} \mid \theta^{(i)})$ .

In practice, we cannot generate iid samples from  $p(\theta \mid y)$ .

Markov chain Monte Carlo:

- Start with an initial draw  $\theta^{(0)} \sim p_0(\theta)$ .
- Apply a transition kernel,  $\theta^{(i+1)} \sim \Gamma(\theta^{(i+1)} \mid \theta^{(i)})$ .

Under certain conditions,

$$\theta^{(n)} \xrightarrow[n \rightarrow \infty]{d} p(\theta \mid y),$$

and  $p(\theta \mid y)$  is the [stationary distribution](#).

In practice, we cannot generate iid samples from  $p(\theta \mid y)$ .

Markov chain Monte Carlo:

- Start with an initial draw  $\theta^{(0)} \sim p_0(\theta)$ .
- Apply a transition kernel,  $\theta^{(i+1)} \sim \Gamma(\theta^{(i+1)} \mid \theta^{(i)})$ .

Under certain conditions,

$$\theta^{(n)} \xrightarrow[n \rightarrow \infty]{d} p(\theta \mid y),$$

and  $p(\theta \mid y)$  is the [stationary distribution](#).

In practice, for “large enough”  $n$ ,

$$\theta^{(n)} \stackrel{\text{approx.}}{\sim} p(\theta \mid y).$$

In practice, we cannot generate iid samples from  $p(\theta \mid y)$ .

Markov chain Monte Carlo:

- Start with an initial draw  $\theta^{(0)} \sim p_0(\theta)$ .
- Apply a transition kernel,  $\theta^{(i+1)} \sim \Gamma(\theta^{(i+1)} \mid \theta^{(i)})$ .

Under certain conditions,

$$\theta^{(n)} \xrightarrow[n \rightarrow \infty]{d} p(\theta \mid y),$$

and  $p(\theta \mid y)$  is the **stationary distribution**.

In practice, for “large enough”  $n$ ,

$$\theta^{(n)} \stackrel{\text{approx.}}{\sim} p(\theta \mid y).$$

- The first samples suffer from a large bias.

In practice, we cannot generate iid samples from  $p(\theta \mid y)$ .

Markov chain Monte Carlo:

- Start with an initial draw  $\theta^{(0)} \sim p_0(\theta)$ .
- Apply a transition kernel,  $\theta^{(i+1)} \sim \Gamma(\theta^{(i+1)} \mid \theta^{(i)})$ .

Under certain conditions,

$$\theta^{(n)} \xrightarrow[n \rightarrow \infty]{d} p(\theta \mid y),$$

and  $p(\theta \mid y)$  is the **stationary distribution**.

In practice, for “large enough”  $n$ ,

$$\theta^{(n)} \stackrel{\text{approx.}}{\sim} p(\theta \mid y).$$

- The first samples suffer from a large bias.
- Discard these samples during a burn-in or *warmup* phase.



**Example: Metropolis algorithm** [Metropolis et al., 1953]

## Example: Metropolis algorithm [Metropolis et al., 1953]

- 1 Start at an initial point,  $\theta^{(0)} \sim p_0$ .

## Example: Metropolis algorithm [Metropolis et al., 1953]

- 1 Start at an initial point,  $\theta^{(0)} \sim p_0$ .
- 2 Apply the transition kernel  $N$  times:

- 1 Propose a new sample

$$\theta^{(i+1)} \sim \text{normal}(\theta^{(i)}, \sigma^2 I)$$

- 2 Accept the proposal with probability

$$\text{Pr} = \min\left(\frac{p(\theta^{(i+1)} | z)}{p(\theta^{(i)} | z)}, 1\right).$$

## Example: Metropolis algorithm [Metropolis et al., 1953]

- 1 Start at an initial point,  $\theta^{(0)} \sim p_0$ .
- 2 Apply the transition kernel  $N$  times:

- 1 Propose a new sample

$$\theta^{(i+1)} \sim \text{normal} \left( \theta^{(i)}, \sigma^2 I \right)$$

- 2 Accept the proposal with probability

$$\text{Pr} = \min \left( \frac{p(\theta^{(i+1)} \mid z)}{p(\theta^{(i)} \mid z)}, 1 \right).$$

- 3 Return the chain  $(\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(N)})$ .

## Example: Metropolis algorithm [[Metropolis et al., 1953](#)]




Figure from [[Gelman et al., 2013](#)].

## Example: Metropolis algorithm [Metropolis et al., 1953]

Benefits:

- Only requires evaluating  $p(\theta, y) = p(\theta)p(y | \theta)$ .
- Asymptotically, the algorithm samples from  $p(\theta | y)$ .

Drawbacks:

- In the finite regime, the samples are **biased**.
- The samples are not independent; there are correlated, which **increases the variance** of our Monte Carlo estimators.

## Example: Continuous diffusion process

In the limit where we take infinitesimally small steps, many MCMC algorithms can be approximated by a random diffusion process

[Gelman et al., 1997, Roberts and Rosenthal, 1998].

- Initial distribution:  $p_0 = \text{normal}(\mu_0, \sigma_0^2)$ .
- Target distribution:  $p = \text{normal}(\mu, \sigma^2)$ .

## Example: Continuous diffusion process

In the limit where we take infinitesimally small steps, many MCMC algorithms can be approximated by a random diffusion process

[Gelman et al., 1997, Roberts and Rosenthal, 1998].

- Initial distribution:  $p_0 = \text{normal}(\mu_0, \sigma_0^2)$ .
- Target distribution:  $p = \text{normal}(\mu, \sigma^2)$ .

Then after time  $T$ ,

$$\theta^{(T)} \sim \text{normal} \left[ (\mu_0 - \mu)e^{-T} + \mu, \quad (\sigma_0^2 - \sigma^2) e^{-2T} + \sigma^2 \right].$$



## Example: Continuous diffusion process

In the limit where we take infinitesimally small steps, many MCMC algorithms can be approximated by a random diffusion process

[Gelman et al., 1997, Roberts and Rosenthal, 1998].

- Initial distribution:  $p_0 = \text{normal}(\mu_0, \sigma_0^2)$ .
- Target distribution:  $p = \text{normal}(\mu, \sigma^2)$ .

Then after time  $T$ ,

$$\theta^{(T)} \sim \text{normal} \left[ (\mu_0 - \mu)e^{-T} + \mu, \quad (\sigma_0^2 - \sigma^2) e^{-2T} + \sigma^2 \right].$$

*For  $T$  large enough, the bias becomes negligible.*

## Variance of Monte Carlo estimator

Suppose the chain is *stationary*; i.e. we started at  $p_0 = p(\theta \mid y)$  or we already ran the chain for an infinitely long time.

## Variance of Monte Carlo estimator

Suppose the chain is *stationary*; i.e. we started at  $p_0 = p(\theta \mid y)$  or we already ran the chain for an infinitely long time.

- Under certain conditions, Monte Carlo estimators observe a central limit theorem, meaning that for large  $N$ ,

$$\frac{1}{N} \sum_i f(\theta^{(n)}) \stackrel{\text{approx}}{\sim} \text{Normal} \left( \mathbb{E}[f(\theta)], \frac{\text{Var} f(\theta)}{N_{\text{eff}}} \right)$$

where  $N_{\text{eff}}$  is the **effective sample size (ESS)**.

## Variance of Monte Carlo estimator

Suppose the chain is *stationary*; i.e. we started at  $p_0 = p(\theta \mid y)$  or we already ran the chain for an infinitely long time.

- Under certain conditions, Monte Carlo estimators observe a central limit theorem, meaning that for large  $N$ ,

$$\frac{1}{N} \sum_i f(\theta^{(n)}) \stackrel{\text{approx}}{\sim} \text{Normal} \left( \mathbb{E}[f(\theta)], \frac{\text{Var} f(\theta)}{N_{\text{eff}}} \right)$$

where  $N_{\text{eff}}$  is the **effective sample size (ESS)**.

- 

$$N_{\text{eff}} = \frac{N}{1 + 2 \sum_{t=1}^{\infty} \rho_t}.$$

$\rho_t$  is the chain's autocorrelation between  $\theta^{(i)}$  and  $\theta^{(i+t)}$ .

## Handling the error of MCMC



In practice, MCMC proceeds in two phases:

## Handling the error of MCMC



In practice, MCMC proceeds in two phases:

**Warmup phase:** We run the process for several steps for the bias to become negligible but don't use any of those samples in our Monte Carlo estimator.

## Handling the error of MCMC



In practice, MCMC proceeds in two phases:

**Warmup phase:** We run the process for several steps for the bias to become negligible but don't use any of those samples in our Monte Carlo estimator.


**Sampling phase:** Collect enough samples to have a large ESS and reduce the variance of the Monte Carlo estimator.

**Question:** Which transition kernel should we choose? Many choices!


Metropolis, Gibbs, Metropolis-adjusted Langevin approximation, Hamiltonian Monte Carlo,...




Example: ill-conditioned Gaussian




Kernel: Metropolis-Hastings proposal (low acceptance probability)



Kernel: Metropolis-Hastings proposal (small step)



Kernel: Gibbs sampler (moves one coordinate at a time)



# Hamiltonian Monte Carlo

- Treat the Markov chain as a physical particle, which evolves over  $\mathbb{R}^D$ , subject to a *potential*:

$$U(\theta) = -\log p(\theta \mid y).$$

# Hamiltonian Monte Carlo

- Treat the Markov chain as a physical particle, which evolves over  $\mathbb{R}^D$ , subject to a *potential*:

$$U(\theta) = -\log p(\theta \mid y).$$

- Give the particle a random shove, by giving it a *momentum*  $\xi_0 \in \mathbb{R}^D$ ,

$$\xi_0 \sim \text{normal}(0, M)$$

# Hamiltonian Monte Carlo

- Treat the Markov chain as a physical particle, which evolves over  $\mathbb{R}^D$ , subject to a *potential*:

$$U(\theta) = -\log p(\theta \mid y).$$


- Give the particle a random shove, by giving it a *momentum*  $\xi_0 \in \mathbb{R}^D$ ,

$$\xi_0 \sim \text{normal}(0, M)$$

- Simulate a the laws of classical mechanics for a time  $T$ ,

$$(\theta_0, \xi_0) \rightarrow (\theta_T, \xi_T).$$

$$\frac{d\theta}{dt} = M^{-1}\xi, \quad \frac{d\xi}{dt} = \nabla_{\theta} \log p(\theta \mid y).$$




A particle accelerates when  $U(\theta) = -\log p(\theta \mid y)$  decreases.

It decelerates when  $U(\theta) = -\log p(\theta \mid y)$  increases.

This is based on  $\nabla_{\theta} \log p(\theta \mid y)$ .



Kernel: Hamiltonian Monte Carlo




Some challenges in implementing HMC:

- Need to compute  $\nabla_{\theta} \log p(\theta | y) \rightarrow$  automatic differentiation
- Tuning parameters:
  - The length of the Hamiltonian trajectory
  - The precision with which we solve Hamilton's equations of motion (step size)
  - The mass matrix of the fictitious particle

Some challenges in implementing HMC:

- Need to compute  $\nabla_{\theta} \log p(\theta | y) \rightarrow$  automatic differentiation
- Tuning parameters:
  - The length of the Hamiltonian trajectory
  - The precision with which we solve Hamilton's equations of motion (step size)
  - The mass matrix of the fictitious particle
- **Stan** implements autodiff and the No-U-Turn Sampler (NUTS) [Hoffman and Gelman, 2014].
  - Adaptive setting of trajectory length to avoid U-Turns
  - Adaptive tuning of step size and mass matrix during warmup



- Figure from [Hoffman and Gelman, 2014].
- Geometric structures are common in complex models and arise naturally in high dimensions.
- **Stan** implements an improved NUTS sampler [Betancourt, 2018]. I highly recommends Betancourt's paper on HMC for those who want to dive deeper!

### III

Application: Bayesian Linear Regression  
(warmup example)

## How Stan works

- The Stan file specifies the joint distribution

$$p(\theta, y) = p(y|\theta)p(\theta) \propto p(\theta | y)$$

- The input includes:
  - the data,  $y$
  - tuning parameters for the algorithm
- The output can include:
  - an approximate sample from the posterior distribution
  - summaries of the run which can help us diagnose problems.

## Inference algorithms in Stan

- Hamiltonian Monte Carlo (HMC)
- No-U Turn Sampler (NUTS)
- Automatic differentiation variational inference (ADVI)
- Pathfinder Variational Inference
- ...

We can manage the **Stan** file, the input, and the output using a scripting language, such as:

- R
- Python
- Julia
- The command line
- . . .

The **Stan** documentation is your friend:

<https://mc-stan.org/users/documentation/!!>

And so is its community: <https://discourse.mc-stan.org/>.



## Example: Bayesian linear regression

The data generating process is:

$$p(\beta) = \text{Normal}(2, 1)$$

$$p(\sigma) = \text{Normal}^+(1, 1)$$

$$p(y \mid \beta, \sigma) = \text{Normal}(\beta x, \sigma)$$

Our goal is to estimate  $\theta = (\beta, \sigma)$ , based on the observation  $z = (x, y)$  and prior knowledge we have of  $\beta$  and  $\sigma$ .

## Writing the Stan file

Stan retains certain C++ features:

- Variables need to be declared.
- Each statement must end with a semi-colon.

For example:

```
real x;
```

## Writing the Stan file

```
data {  
  Declare the data that will be given as an input.  
}
```

```
parameters {  
  Declare the parameters we want to sample.  
}
```

```
model {  
  Compute the log joint distribution.  
}
```

[code demo](#)

## Convergence diagnostic

Are the chains still biased by their initializations?

**Proposition:** Start multiple chains at a different locations and check that they all converge to the same distribution.

Examine:

- 1 the trace plots
- 2 the density plots
- 3 the  $\hat{R}$  statistic

$$\hat{R} := \frac{\text{Standard deviation across all chains}}{\text{Standard deviation within chain}}$$

- If the chains sample from the same target, expect  $\hat{R} \approx 1$ .
- If the chains are disagreement,  $\hat{R} \gg 1$ .

code demo

## A more in-depth look at $\hat{R}$

Let  $\theta^{(nm)}$  be the  $n^{\text{th}}$  sample from the  $m^{\text{th}}$  chain.

Let  $f : \Theta \rightarrow \mathbb{R}$  be some function of interest.

## A more in-depth look at $\hat{R}$

Let  $\theta^{(nm)}$  be the  $n^{\text{th}}$  sample from the  $m^{\text{th}}$  chain.

Let  $f : \Theta \rightarrow \mathbb{R}$  be some function of interest.

Can write  $\hat{R}$  as

$$\hat{R} = \sqrt{\frac{N-1}{N} + \frac{\hat{B}}{\widehat{W}}},$$

where

- $\hat{B}$  is the sample variance of  $\bar{f}(\theta^{(\cdot m)})$ .
- $\widehat{W}$  is the (average) within-chain variance.

## A more in-depth look at $\hat{R}$

Let  $\theta^{(nm)}$  be the  $n^{\text{th}}$  sample from the  $m^{\text{th}}$  chain.

Let  $f : \Theta \rightarrow \mathbb{R}$  be some function of interest.

Can write  $\hat{R}$  as

$$\hat{R} = \sqrt{\frac{N-1}{N} + \frac{\hat{B}}{\widehat{W}}},$$

where

- $\hat{B}$  is the sample variance of  $\bar{f}(\theta^{(\cdot m)})$ .
- $\widehat{W}$  is the (average) within-chain variance.

$$\hat{R} \leq 1 + \epsilon \iff \hat{B} \lesssim 2\epsilon\widehat{W} + \mathcal{O}(\epsilon^2).$$

Want to make sure  $\text{Var}[f(\bar{\theta}^{(\cdot m)})]$  is small.

A more in-depth look at  $\hat{R}$



**Question.** What can  $\text{Var} [f(\bar{\theta}^{(\cdot m)})]$  teach us about convergence and bias decay?



A more in-depth look at  $\hat{R}$



**Question.** What can  $\text{Var} [f(\bar{\theta}^{(\cdot m)})]$  teach us about convergence and bias decay?

$$\text{Var} \left( \bar{f}^{(\cdot m)} \right) = \text{Var} \left( \mathbb{E} \left( \bar{f}^{(\cdot m)} \mid f^{(0)} \right) \right) + \mathbb{E} \left( \text{Var} \left( \bar{f}^{(\cdot m)} \mid f^{(0)} \right) \right).$$

## A more in-depth look at $\hat{R}$



**Question.** What can  $\text{Var} [f(\bar{\theta}^{(\cdot m)})]$  teach us about convergence and bias decay?

$$\text{Var} \left( \bar{f}^{(\cdot m)} \right) = \text{Var} \left( \mathbb{E} \left( \bar{f}^{(\cdot m)} \mid f^{(0)} \right) \right) + \mathbb{E} \left( \text{Var} \left( \bar{f}^{(\cdot m)} \mid f^{(0)} \right) \right).$$

The **nonstationary variance** measures how well the chains forget their starting points.

## A more in-depth look at $\hat{R}$



**Question.** What can  $\text{Var} [f(\bar{\theta}^{(\cdot m)})]$  teach us about convergence and bias decay?

$$\text{Var} \left( \bar{f}^{(\cdot m)} \right) = \text{Var} \left( \mathbb{E} \left( \bar{f}^{(\cdot m)} \mid f^{(0)} \right) \right) + \mathbb{E} \left( \text{Var} \left( \bar{f}^{(\cdot m)} \mid f^{(0)} \right) \right).$$

The **nonstationary variance** measures how well the chains forget their starting points.

As we warmup the chains, both the **nonstationary variance** and **squared bias** decay to 0, and so  $\hat{R}$  acts as a “proxy clock” for bias.

## A more in-depth look at $\hat{R}$

- What quantity does  $\hat{R}$  measure and how close to 1 should it be?
  - [Vehtari et al., 2021] propose checking that  $\hat{R} \leq 1.01$ .
  - [Margossian et al., 2024] examine  $\hat{R}$  for nonstationary chains and propose a more direct measure of the nonstationary variance.



- Ideally,  $\hat{R}$  tells us if the warmup phase is long enough.  
In practice, it tells us if both the warmup and sampling phase are long enough.



- Ideally,  $\hat{R}$  tells us if the warmup phase is long enough.  
In practice, it tells us if both the warmup and sampling phase are long enough.
- ESS and Monte Carlo standard error (MCSE) tell us if the sampling phase is long enough.  
Rule of thumb: aim for  $\text{ESS} \geq 100$ . (we'll think harder about that)



- Ideally,  $\hat{R}$  tells us if the warmup phase is long enough.  
In practice, it tells us if both the warmup and sampling phase are long enough.
- ESS and Monte Carlo standard error (MCSE) tell us if the sampling phase is long enough.  
Rule of thumb: aim for  $\text{ESS} \geq 100$ . (we'll think harder about that)
- $\text{ESS}_{\text{tail}}$  quantifies information for tail estimates [Vehtari et al., 2021].



- Ideally,  $\hat{R}$  tells us if the warmup phase is long enough.  
In practice, it tells us if both the warmup and sampling phase are long enough.
- ESS and Monte Carlo standard error (MCSE) tell us if the sampling phase is long enough.  
Rule of thumb: aim for  $\text{ESS} \geq 100$ . (we'll think harder about that)
- $\text{ESS}_{\text{tail}}$  quantifies information for tail estimates [Vehtari et al., 2021].
- Median,  $M(\theta)$  and Median Absolute Deviation (MAD),


$$M(|\theta^{(i)} - M(\theta)|)$$

can be helpful when the first moments are not finite.



## Posterior predictive checks

- Recall Box's loop (from *Bayesian Workflow*)!
- Does our model accurately describe the data?



## Posterior predictive checks (“check trained model”)

### Proposition:

Each time we draw a sample,  $\theta^{(i)} = (\beta^{(i)}, \sigma^{(i)})$ , we will also simulate data, according to:

$$y_{\text{pred}}^{(i)} \sim \text{Normal} \left( x\beta^{(i)}, \sigma^{(i)} \right)$$

## Posterior predictive checks (“check trained model”)

### Proposition:

Each time we draw a sample,  $\theta^{(i)} = (\beta^{(i)}, \sigma^{(i)})$ , we will also simulate data, according to:

$$y_{\text{pred}}^{(i)} \sim \text{Normal} \left( x\beta^{(i)}, \sigma^{(i)} \right)$$

Want to study the posterior predictive distribution,

$$p(y_{\text{pred}} \mid y) = \int_{\Theta} p(y_{\text{pred}} \mid \theta) p(\theta \mid y) d\theta.$$

code demo

## Improving the model


- The ppc suggest our model can improve with an intercept parameter.
- *Exercise:* repeat the above procedure, but this time add an intercept parameter  $\beta_0$ . Check that the inference is reliable and perform new posterior predictive checks.

### III

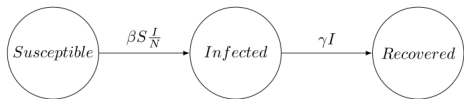
Disease transmission model

1978 influenza outbreak in a British boarding school.

Data: daily number of students in bed.



## Susceptible-Infected-Recovered (SIR) model




$N$ : total number of individuals,  
 $N = S + I + R$ .

$$\dot{S} = -\beta SI/N$$

$$\dot{I} = \beta SI/N - \gamma I$$

$$\dot{R} = \gamma I$$

## Susceptible-Infected-Recovered (SIR) model



$$\dot{S} = -\beta SI/N$$

$$\dot{I} = \beta SI/N - \gamma I$$

$$\dot{R} = \gamma I$$


$N$ : total number of individuals,  
 $N = S + I + R$ .

$\beta$ : transmission rate.

$\gamma$ : rate of recovery of infected individuals.



## Susceptible-Infected-Recovered (SIR) model



$$\dot{S} = -\beta SI/N$$

$$\dot{I} = \beta SI/N - \gamma I$$

$$\dot{R} = \gamma I$$

$N$ : total number of individuals,  
 $N = S + I + R$ .

$\beta$ : transmission rate.

$\gamma$ : rate of recovery of infected individuals.

- $I/N$ : the proportion of infectious individuals.
- $\beta(I/N)$ : the probability that a single susceptible individual becomes infected in one day.

Which measurement model should we use?

Which measurement model should we use?

- ④ *Poisson* likelihood parameterized by  $\lambda(t) = I(t)$ .
  - Then  $\mathbb{E}(y(t)) = I(t)$  and  $\text{Var}(y(t)) = I(t)$ .

Which measurement model should we use?

- ① *Poisson* likelihood parameterized by  $\lambda(t) = I(t)$ .
  - Then  $\mathbb{E}(y(t)) = I(t)$  and  $\text{Var}(y(t)) = I(t)$ .
  
- ② *Negative-Binomial* parameterized by  $\mu = I(t)$  and  $\phi$ .
  - Then  $\mathbb{E}(y(t)) = I(t)$  and  $\text{Var}(y(t)) = I(t) + \frac{I(t)^2}{\phi}$ .

Which measurement model should we use?

- ① *Poisson* likelihood parameterized by  $\lambda(t) = I(t)$ .
  - Then  $\mathbb{E}(y(t)) = I(t)$  and  $\text{Var}(y(t)) = I(t)$ .
- ② *Negative-Binomial* parameterized by  $\mu = I(t)$  and  $\phi$ .
  - Then  $\mathbb{E}(y(t)) = I(t)$  and  $\text{Var}(y(t)) = I(t) + \frac{I(t)^2}{\phi}$ .
  - In **Stan** use `neg_binomial_2`.
  - Define in `parameters` block  $\phi^{-1}$ .

Which prior should we use?

- $p(\beta) = \text{normal}^+(2, 1)$ : restricts  $\beta$  to be positive and  $p(\beta < 4) = 0.975$ .
- $p(\gamma) = \text{normal}^+(0.4, 0.5)$ : restricts  $\gamma$  to be positive and  $p(\gamma < 1) = 0.9$ , i.e. 90% of the time, we expect the average time spent in bed to be less than 1 day).
- $p(\phi^{-1}) = \text{exponential}(5)$ , see [Grinsztajn et al., 2021].

code demo

Additional controls for the MCMC sampler.

- Number of (parallel) chains
- Number of warmup iterations
- Number of sampling iterations

*Exercise: Write and fit an SIR model for the 1978 influenza outbreak.*

- *Check the standard diagnostics ( $\hat{R}$  and ESS) and examine the density and trace plots. Is the inference reliable?*
- *Do posterior predictive checks: does the model accurately describe the data?*
- *Report  $\beta$ ,  $\gamma$ ,  $R_0 = \beta/\gamma$  and the recovery time  $T = 1/\gamma$ .*
- *How precise are your estimates of  $T$ ? Should you run longer Markov chains? Shorter ones?*
- *Compare the two proposed measurement models: Poisson and negative binomial.*

For more discussion about this model (e.g. choice of priors, sensitivity tests), see [Grinsztajn et al., 2021].

For more discussion about the length of Markov chains, see [?].



# IV

## Model Comparison

**Question:** For the SIR model, do we get better predictions with the Poisson or the negative binomial likelihood?

**Question:** For the SIR model, do we get better predictions with the Poisson or the negative binomial likelihood?

- **Proposition:** Test *model predictions* on a validation set.

**Question:** For the SIR model, do we get better predictions with the Poisson or the negative binomial likelihood?

- **Proposition:** Test *model predictions* on a validation set.
  - Split the data into a **training** and a **validation** set.

**Question:** For the SIR model, do we get better predictions with the Poisson or the negative binomial likelihood?

- **Proposition:** Test *model predictions* on a validation set.
  - Split the data into a **training** and a **validation** set.
- **Training set:** The data  $y_{\text{train}}$  used to learn the parameters, and on which we condition the posterior,

$$p(\theta \mid y_{\text{train}}).$$

**Question:** For the SIR model, do we get better predictions with the Poisson or the negative binomial likelihood?

- **Proposition:** Test *model predictions* on a validation set.

- Split the data into a **training** and a **validation** set.

- **Training set:** The data  $y_{\text{train}}$  used to learn the parameters, and on which we condition the posterior,

$$p(\theta \mid y_{\text{train}}).$$

- **Validation set:** The data  $y_{\text{val}}$  we use to “test” the model’s predictions.

Example: At  $t = 12$ , the model predicts  $\tilde{y}(t = 12)$ . Compute the *prediction error*,

$$\text{Err} = (\tilde{y}(t = 12) - y_{\text{val}}(t = 12))^2.$$

Testing *uncertainty calibration* in (point) predictions

Testing *uncertainty calibration* in (point) predictions

Suppose we have a normal likelihood, with point estimates for the learned parameters,

$$\text{Normal}(\hat{\mu}(t), \hat{\sigma}).$$

Our “best” prediction is  $\tilde{y}(t) = \hat{\mu}(t)$ .



Testing *uncertainty calibration* in (point) predictions

Suppose we have a normal likelihood, with point estimates for the learned parameters,

$$\text{Normal}(\hat{\mu}(t), \hat{\sigma}).$$

Our “best” prediction is  $\tilde{y}(t) = \hat{\mu}(t)$ .

Then the prediction error is

$$\text{Err} = (\hat{\mu}(t) - y_{\text{val}}(t))^2,$$

and  $\hat{\sigma}$  is unaccounted for!

Testing *uncertainty calibration* in (point) predictions

Suppose we have a normal likelihood, with point estimates for the learned parameters,

$$\text{Normal}(\hat{\mu}(t), \hat{\sigma}).$$

Our “best” prediction is  $\tilde{y}(t) = \hat{\mu}(t)$ .

Then the prediction error is

$$\text{Err} = (\hat{\mu}(t) - y_{\text{val}}(t))^2,$$

and  $\hat{\sigma}$  is unaccounted for!

Instead, let's evaluate the *point-estimate log predictive density*,

$$\begin{aligned} \text{p-lpd} &= \log p(y_{\text{val}}(t) \mid \hat{\mu}, \hat{\sigma}) \\ &= \text{const.} - \log \hat{\sigma} - \frac{1}{2\hat{\sigma}^2} (y_{\text{val}}(t) - \hat{\mu}(t))^2. \end{aligned}$$

Testing *uncertainty calibration* in (point) predictions

Suppose we have a Bernoulli likelihood, with point estimates for the learned parameters,

$$\text{Bernoulli}(\hat{\pi}(t)).$$

Our “best” prediction is  $\tilde{y}(t) = \mathbb{I}[\hat{\pi}(t) > 0.5]$ .

Then the prediction error is

$$\text{Err} = \mathbb{I}[\tilde{y}(t) = y_{\text{val}}(t)].$$

Instead, let's evaluate the *point-estimate log predictive density*,

$$\begin{aligned} \text{p-lpd} &= \log p(y_{\text{val}}(t) \mid \hat{\pi}(t)) \\ &= y_{\text{val}}(t) \log \hat{\pi}(t) + (1 - y_{\text{val}}(t)) \log(1 - \hat{\pi}(t)). \end{aligned}$$

Testing *uncertainty calibration* in Bayesian predictions

We have a general strategy which accounts for uncertainty in the likelihood for a fixed  $\theta$ ,

$$\text{p-lpd} = \log p(y_{\text{val}}(t) \mid \theta).$$

Testing *uncertainty calibration* in Bayesian predictions

We have a general strategy which accounts for uncertainty in the likelihood for a fixed  $\theta$ ,

$$\text{p-lpd} = \log p(y_{\text{val}}(t) \mid \theta).$$

In a Bayesian framework, we integrate with respect to the posterior and obtain the *expected log predictive density*,

$$\begin{aligned} \text{elpd} &= \log p(y_{\text{val}}(t) \mid y_{\text{train}}) \\ &= \log \int_{\Theta} p(y_{\text{val}}(t) \mid \theta) p(\theta \mid y_{\text{train}}) d\theta. \end{aligned}$$

How do we split the data  $(t, y)$  into a training and a test set?

**Proposition:** Do *leave-one-out cross validation* and compute

$$\text{elpd}_{\text{loo}} = \sum_{i=1}^N \log p(y_i \mid y_{-i}),$$

where

$$p(y_i \mid y_{-i}) = \int_{\Theta} p(y_i \mid \theta) p(\theta \mid y_{-i}) d\theta.$$

## Recap.

Prediction error based on “best” prediction,  $(y_{\text{val}} - \tilde{y})^2$

→ point-wise log predictive score,  $\text{p-lpd} = \log p(y_{\text{val}} \mid \hat{\theta})$

→ expected log predictive score,  $\text{elpd} = \log p(y_{\text{val}} \mid y_{\text{train}})$

→ loo CV,  $\text{elpd}_{\text{loo}} = \sum_{i=1}^N \log p(y_i \mid y_{-i})$

How do we estimate  $\text{elpd}_{\text{loo}}$  efficiently?

## Importance Sampling

- **Idea:** Suppose we need to estimate an expectation with respect to  $\ell(\theta)$ ,

$$\int_{\Theta} f(\theta) \ell(\theta) d\theta,$$

but using samples from  $q(\theta)$ .



## Importance Sampling

- **Idea:** Suppose we need to estimate an expectation with respect to  $\ell(\theta)$ ,

$$\int_{\Theta} f(\theta) \ell(\theta) d\theta,$$

but using samples from  $q(\theta)$ .

- For example,  $\ell(\theta) = p(\theta \mid y_{-i})$  and  $q(\theta) = p(\theta \mid y)$ .

## Importance Sampling

- **Idea:** Suppose we need to estimate an expectation with respect to  $\ell(\theta)$ ,

$$\int_{\Theta} f(\theta) \ell(\theta) d\theta,$$

but using samples from  $q(\theta)$ .

- For example,  $\ell(\theta) = p(\theta \mid y_{-i})$  and  $q(\theta) = p(\theta \mid y)$ .
- Then, note that

$$\int_{\Theta} f(\theta) \ell(\theta) d\theta = \int_{\Theta} f(\theta) \frac{\ell(\theta)}{q(\theta)} q(\theta) d\theta.$$

## Importance Sampling

- **Idea:** Suppose we need to estimate an expectation with respect to  $\ell(\theta)$ ,

$$\int_{\Theta} f(\theta) \ell(\theta) d\theta,$$

but using samples from  $q(\theta)$ .

- For example,  $\ell(\theta) = p(\theta \mid y_{-i})$  and  $q(\theta) = p(\theta \mid y)$ .
- Then, note that

$$\int_{\Theta} f(\theta) \ell(\theta) d\theta = \int_{\Theta} f(\theta) \frac{\ell(\theta)}{q(\theta)} q(\theta) d\theta.$$

- The IS Monte Carlo estimator is

$$\widehat{\mathbb{E}} f(\theta) = \frac{1}{S} \sum_{s=1}^S f(\theta^{(s)}) \frac{\ell(\theta^{(s)})}{q(\theta^{(s)})}.$$

- The IS Monte Carlo estimator is

$$\widehat{\mathbb{E}}f(\theta) = \sum_{s=1}^S f(\theta^{(s)}) \frac{\ell(\theta^{(s)})}{q(\theta^{(s)})}.$$

- The IS Monte Carlo estimator is

$$\hat{\mathbb{E}}f(\theta) = \sum_{s=1}^S f(\theta^{(s)}) \frac{\ell(\theta^{(s)})}{q(\theta^{(s)})}.$$

Practical concern: it is not uncommon for the IS estimator to have a non-finite variance and so we need  $q(\theta) \approx \ell(\theta)$ !

- The IS Monte Carlo estimator is

$$\widehat{\mathbb{E}}f(\theta) = \sum_{s=1}^S f(\theta^{(s)}) \frac{\ell(\theta^{(s)})}{q(\theta^{(s)})}.$$

Practical concern: it is not uncommon for the IS estimator to have a non-finite variance and so we need  $q(\theta) \approx \ell(\theta)$ !

### Proposition

*When the  $y_j$ 's are independent conditioned on  $\theta$ , the importance sampling Monte Carlo estimator is*

$$\widehat{p}(y_i \mid y_{-i}) = \frac{1}{\sum_{s=1}^S \frac{1}{p(y_i \mid \theta^{(s)})}},$$

*where  $\theta^{(s)} \sim p(\theta \mid y)$ .*

Practical concerns:

- Several steps need to be taken to stabilize IS estimators.

Practical concerns:

- Several steps need to be taken to stabilize IS estimators.
- We will use *Pareto-smoothed importance sampling* (PSIS) [Vehtari et al., 2017].



Practical concerns:

- Several steps need to be taken to stabilize IS estimators.
- We will use *Pareto-smoothed importance sampling* (PSIS) [Vehtari et al., 2017].
- PSIS comes equipped with a  $\hat{k}$  diagnostic:
  - if  $\hat{k} < 0.5$ , PSIS estimators is reliable.
  - if  $\hat{k} \geq 0.7$ , importance weights have non-finite variance.

## Practical concerns:

- Several steps need to be taken to stabilize IS estimators.
- We will use *Pareto-smoothed importance sampling* (PSIS) [Vehtari et al., 2017].
- PSIS comes equipped with a  $\hat{k}$  diagnostic:
  - if  $\hat{k} < 0.5$ , PSIS estimators is reliable.
  - if  $\hat{k} \geq 0.7$ , importance weights have non-finite variance.
- The R package `loo` computes PSIS.
- In **Stan** 's generated quantities, need to compute `log_lik`, where
$$\text{log\_lik}[i] = \log p(\text{cases}[i] \mid \text{theta});$$

*Exercise:* Compare the predictive scores of the SIR models.

- Evaluate in **generated quantities** the log probability mass functions using `poisson_lpmf` and `neg_binomial_2_lpmf`.
- In R, use the `loo` package to compute the PSIS estimates of the  $\text{elpd}_{\text{loo}}$ .
- Check  $\hat{k}$  to see if the IS estimators are reliable.
- Which likelihood achieves the best predictive score?

## V

### Concluding Remarks

Where does **Stan** fit in the Bayesian modeler's toolkit?

Historical contribution:

- **Stan** was born around 2012.
- First intended as a well programmed version of BUGS and JAGS.

Where does **Stan** fit in the Bayesian modeler's toolkit?

Historical contribution:

- **Stan** was born around 2012.
- First intended as a well programmed version of BUGS and JAGS.

Several algorithms were developed as part of **Stan** 's development:

- Adaptive Hamiltonian Monte Carlo  
[Hoffman and Gelman, 2014, Betancourt, 2018]
- ADVI: a black box variational inference [Kucukelbir et al., 2017]
- PathFinder: an improved variational inference [Zhang et al., 2022].
- Delayed rejection HMC [Modi et al., 2023]
- Adjoint-differentiated Laplace approximation [Margossian et al., 2020]
- GIST sampler [Bou-Rabee et al., 2024]

**Stan** by the people, for the people

- **Stan** is open source: <https://github.com/stan-dev>
- Contributing new functions to **Stan** : <https://github.com/stan-dev/stan/wiki/Contributing-New-Functions-to-Stan>
- Straightforward to interface **Stan** with your own gradient-based algorithm (written in your language of choice), via BridgeStan:  
<https://roualdes.github.io/bridgestan/latest/>.
- A lot of other probabilistic programming languages out there: PyMC, Turing, TensorFlow Probability, PyTorch, etc.

Where do Monte Carlo methods fit in the probabilistic modeling toolkit?  
(a high-level comparison)

**Markov chain Monte Carlo**

**Variational inference**



Where do Monte Carlo methods fit in the probabilistic modeling toolkit?  
(a high-level comparison)

### **Markov chain Monte Carlo**

- use general transition kernel

### **Variational inference**

- requires an approximating model

Where do Monte Carlo methods fit in the probabilistic modeling toolkit?  
(a high-level comparison)

### **Markov chain Monte Carlo**

- use general transition kernel
- initially error goes down slowly

### **Variational inference**

- requires an approximating model
- initially error goes down quickly

Where do Monte Carlo methods fit in the probabilistic modeling toolkit?  
(a high-level comparison)

### **Markov chain Monte Carlo**

- use general transition kernel
- initially error goes down slowly
- asymptotically unbiased

### **Variational inference**

- requires an approximating model
- initially error goes down quickly
- asymptotically biased

Where do Monte Carlo methods fit in the probabilistic modeling toolkit?  
(a high-level comparison)

### **Markov chain Monte Carlo**

- use general transition kernel
- initially error goes down slowly
- asymptotically unbiased
- several diagnostics available

### **Variational inference**

- requires an approximating model
- initially error goes down quickly
- asymptotically biased
- not obvious how to check inference  
(though good candidates exist)

Where do Monte Carlo methods fit in the probabilistic modeling toolkit?  
(a high-level comparison)

### **Markov chain Monte Carlo**

- use general transition kernel
- initially error goes down slowly
- asymptotically unbiased
- several diagnostics available
  
- can check trained model

### **Variational inference**

- requires an approximating model
- initially error goes down quickly
- asymptotically biased
- not obvious how to check inference  
(though good candidates exist)
  
- can check trained model

Where does MCMC fit in the probabilistic modeling toolkit?




Figure from [Gelman et al., 2020].

## References I

- [Betancourt, 2018] Betancourt, M. (2018).  
A conceptual introduction to Hamiltonian Monte Carlo.  
*arXiv:1701.02434v1*.
- [Bou-Rabee et al., 2024] Bou-Rabee, N., Carpenter, B., and Marsden, M. (2024).  
Gist: Gibbs self-tuning for locally adaptive hamiltonian monte carlo.  
*arXiv 2404.15253*.
- [Gelman et al., 2013] Gelman, A., Carlin, J. B., Stern, H. S., Dunson, D. B., Vehtari, A., and Rubin, D. B. (2013).  
*Bayesian Data Analysis*.  
Chapman & Hall.
- [Gelman et al., 1997] Gelman, A., Gilks, W. R., and Roberts, G. O. (1997).  
Weak convergence and optimal scaling of random walk Metropolis algorithms.  
*Annals of Applied Probability*, 7(1):110–120.
- [Gelman et al., 2020] Gelman, A., Vehtari, A., Simpson, D., Margossian, C. C., Carpenter, B., Yao, Y., Kennedy, L., Gabry, J., Bürkner, P.-C., and Modrák, M. (2020).  
Bayesian workflow.  
*arXiv:2011.01808*.

## References II

- [Grinsztajn et al., 2021] Grinsztajn, L., Semenova, E., Margossian, C. C., and Riou, J. (2021). Bayesian workflow for disease transmission modeling in Stan. *Statistics in Medicine*.
- [Hoffman and Gelman, 2014] Hoffman, M. D. and Gelman, A. (2014). The No-U-Turn Sampler: Adaptively setting path lengths in Hamiltonian Monte Carlo. *Journal of Machine Learning Research*, pages 1593–1623.
- [Kucukelbir et al., 2017] Kucukelbir, A., Tran, D., Ranganath, R., Gelman, A., and Blei, D. (2017). Automatic differentiation variational inference. *Journal of machine learning research*, 18:1 – 45.
- [Margossian et al., 2024] Margossian, C. C., Hoffman, M. D., Sountsov, P., Riou-Durand, L., Vehtari, A., and Gelman, A. (2024). Nested  $\hat{R}$ : Assessing the convergence of Markov chain Monte Carlo when running many short chains. *Bayesian Analysis*.
- [Margossian et al., 2020] Margossian, C. C., Vehtari, A., Simpson, D., and Agrawal, R. (2020). Hamiltonian Monte Carlo using an adjoint-differentiated Laplace approximation: Bayesian inference for latent Gaussian models and beyond. *Neural Information Processing Systems*.



## References III

- [Metropolis et al., 1953] Metropolis, N., Rosenbluth, A., Rosenbluth, M., Teller, A., and Teller, E. (1953). Equations of state calculations by fast computing machines. *Journal of Chemical Physics*, 26.
- [Modi et al., 2023] Modi, C., Barnett, A., and Carpenter, B. (2023). Delayed rejection Hamiltonian Monte Carlo for sampling multiscale distributions. *Bayesian Analysis*.
- [Roberts and Rosenthal, 1998] Roberts, G. O. and Rosenthal, J. S. (1998). Optimal scaling of discrete approximations to Langevin diffusions. *Journal of the Royal Statistical Society, Series B*, 60:255–268.
- [Vehtari et al., 2017] Vehtari, A., Gelman, A., and Gabry, J. (2017). Practical bayesian model evaluation using leave-one-out cross-validation and waic. *Statistics and Computing*, 27:1413–1432.
- [Vehtari et al., 2021] Vehtari, A., Gelman, A., Simpson, D., Carpenter, B., and Bürkner, P.-C. (2021). Rank-normalization, folding, and localization: An improved  $\hat{R}$  for assessing convergence of MCMC (with discussion). *Bayesian Analysis*, 16:667–718.

## References IV

- [Zhang et al., 2022] Zhang, L., Carpenter, B., Gelman, A., and Vehtari, A. (2022).  
Pathfinder: Parallel quasi-Newton variational inference.  
*Journal of Machine Learning Research*, 23(306):1–49.